Subject Review

COMP90042

Natural Language Processing

Lecture 23

Semester 1 2021 Week 12

Jey Han Lau

THE UNIVERSITY OF MELBOURNE

Preprocessing

- Sentence segmentation
- Tokenisation
  - Subword tokenisation
- Word normalisation
  - Derivational vs. inflectional morphology
  - Lemmatisation vs. stemming
- Stop words

N-gram Language Models

- Derivation
- Smoothing techniques
  - Add-*k*
  - Absolute discounting
  - Katz Backoff
  - Kneser-Ney smoothing
  - Interpolation

Text Classification

- Building a classification system
- Text classification tasks
  - Topic classification
  - Sentiment analysis
  - Native language identification
- Algorithms
  - Naive-Bayes, logistic regression, SVM
  - kNN, neural networks
- Bias vs. variance
- Evaluation metrics: Precision, recall, F1

Part-of-Speech Tagging

- English POS
  - Open vs. closed POS classes
- Tagsets
  - Penn Treebank tags
- Automatic taggers
  - Rule-based
  - Statistical
    - Unigram, classifier-based, HMM

Hidden Markov Models

- Probabilistic formulation
  - Parameters: emission and transition probabilities
- Training
- Viterbi algorithm
- Generative vs. discriminative models

DL: Feed-forward Networks

- Formulation
- Designing FF networks for NLP tasks
  - Topic classification
  - Language model
  - POS tagging
- Word embeddings
- Convolutional networks

DL: Recurrent Networks

- Formulation
- RNN language models
- LSTM
  - Functions of gates
  - Variants
- Designing RNN for NLP tasks
  - Text classification: sentiment analysis
  - POS tagging

Lexical Semantics

- Definition of word senses, glosses
- Lexical relationships
  - Synonymy, antonymy, hypernymy, meronymy
- Structure of WordNet
- Word similarity
  - Path length, depth information, information content
- Word sense disambiguation
  - Supervised vs. unsupervised

Distributional Semantics

- Matrices for distributional semantics
  - VSM, TF-IDF, word-word co-occurrence
- Association measures: PMI, PPMI
- Count-based methods: SVD
- Neural methods: skip-gram, CBOW
- Evaluation
  - Word similarity, analogy

Contextual Representation

- Formulation with RNN
- ELMo
- BERT
  - Objectives
  - Fine-tuning for downstream tasks
- Transformers
  - Multi-head attention

Discourse

- Motivation for modelling beyond words
- Discourse segmentation
  - Text Tiling
- Discourse parsing
  - Rhetorical structure theory
- Anaphora resolution
  - Centering
  - Supervised models

Formal Language Theory & FSA

- Formal language theory as a framework for defining language
- Regular languages
  - Closure properties
- Finite state acceptors
  - Word morphology, weighted variant
- Finite state transducers
  - Weighted variant, edit distance, morphological analysis

Context-Free Grammar

- Center embedding
- Basics of CFG
- Syntactic constituent and its properties
- CFG parsing
  - Chomsky normal form
  - CYK
- English sentence structure (Penn Treebank)

Probabilistic Context-Free Grammar

- Ambiguity in grammars
- Basics of probabilistic CFGs
- Probability of a CFG tree
- Parsing
  - Probabilistic CYK
- Improvements
  - Parent conditioning
  - Head lexicalisation

Dependency Grammar

- Notion of dependency between words
- Universal dependency
- Properties of dependency trees
  - Projectivity
- Parsing
  - Transition-based
  - Graph-based

Machine Translation

- Statistical MT
  - Language + translation model
  - Alignments
- Neural MT
  - Encoder-decoder
  - Beam search decoding
  - Attention mechanism
- Evaluation: BLEU

Information Extraction

- Named entity recognition
  - NER tags, IOB tagging, models
- Relation extraction
  - Rule-based, supervised, semi-supervised, distant supervision
  - Unsupervised
- Temporal expression extraction
- Event extraction

Question Answering

- IR-based QA
  - Question processing, answer type prediction
  - Passage retrieval, answer extraction
- Reading comprehension
  - Models: LSTM-based, BERT
- Knowledge-based QA
- Hybrid QA: IBM Watson

Topic Modelling

- Evolution of topic models
- LDA
  - Sampling-based learning
  - Hyper-parameters
- Evaluation:
  - Word intrusion
  - Topic coherence

Summarisation

- Extractive summarisation
  - Single-document
    - Unsupervised content selection
  - Multi-document
    - Maximum marginal relevance
- Abstractive summarisation
  - Encoder-decoder models: copy mechanism
- Evaluation: ROUGE

Ethics

- Learning outcomes
- Arguments against ethical checks
- Core NLP ethics concepts:
  - bias, dual use, privacy
- Discussion of applications/use cases

Exam

Exam Structure

- Open book
- 120 points (40% for the subject)
- Gradescope
- Time: 120 minutes writing time +15 minutes reading time
- 3 parts:
  - A: short answer questions
  - B: method questions
  - C: algorithm questions

Short Answer Questions

- Several short questions
  - 1-2 sentence answers for each
  - Definitional, e.g. what is X?
  - Conceptual, e.g. relate X and Y, purpose of Z?
  - May call for an example illustrating a technique/ problem

Method Questions

- Longer answer
- Focus on analysis and understanding
  - Contrast different methods
  - Analyse an algorithm/application
  - Motivate a modelling technique
  - Explain or derive mathematical equation

Algorithmic Questions

- Perform algorithmic computations
  - Numerical computations for algorithm on some given example data
  - Present an outline of an algorithm on your own example
- Not required to simplify maths (e.g. leaving fractions as log(5/4) is fine)

What to Expect

- Even coverage of topic from the semester
- Be prepared for concepts that have not yet been assessed by homework / project
- Prescribed reading is *fair game* for topics mentioned in the lectures and workshops
- Mock exam

Questions?

- Final survey: https://forms.gle/CYBfYuEh46mjGqG86